

Flight Data Analytics

Overview

Flight Data Analytics analyzes aviation accident data to identify potential risks associated with expansion to the aviation industry as well as advise on how to mitigate said risks. The analysis leverages accident data to inform decisions on aircraft acquisition, equipment purchase as well as staff training to successfully purchase and operate planes for commercial and private enterprises.

Business Problem

Our company is expanding into aviation and needs to evaluate potential risks associated with managing fleets of airplanes. The goal is to: 1.determine which aircraft are best to purchase 2.Necessary equipment to be purchased to enhance safety 3.Necessary experience needed by staff as well as additional training they should undertake

Data Understanding

The data is sourced from the National Transportation Safety Board (NTSB), covering aviation accidents from 1948 to the end of 2022. It includes information such as: 1.accident severity 2.aircraft type 3.injury statistics 3.weather condition 4.plane make and model

For this particular project we will only analyse data going back six(6) years for actionable insight as requested by management

In [28]: ▶ *#importing standard libraries*

```
from matplotlib import pyplot as plt
import pandas as pd
import seaborn as sns
```

In [29]: ▶ `pd.set_option("display.max_columns",500)` *#This allows me to look at all*
`data=pd.read_csv(r"C:\Users\User\Desktop\DATA SCIENCE\Phase 1 project\A`

```
In [30]: df=data.copy()#Create a copy of the data to use for manipulation
df.head() #Check the first five rows make sure data is successfully impo
```

Out[30]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Co
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	L
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	L
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	L
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	L
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	L

```
In [31]: df.shape #Checking the number of rows and columns
```

Out[31]: (88889, 31)

```
In [32]: df.columns #Check for column names
```

Out[32]: Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',
'Airport.Name', 'Injury.Severity', 'Aircraft.damage',
'Aircraft.Category', 'Registration.Number', 'Make', 'Model',
'Amateur.Built', 'Number.ofEngines', 'Engine.Type', 'FAR.Description',
'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',
'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',
'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
'Publication.Date'],
dtype='object')

In [33]: `df.info()` *#Summary of the data*

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                    88889 non-null  object
2   Accident.Number                      88889 non-null  object
3   Event.Date                           88889 non-null  object
4   Location                             88837 non-null  object
5   Country                             88663 non-null  object
6   Latitude                             34382 non-null  object
7   Longitude                            34373 non-null  object
8   Airport.Code                         50132 non-null  object
9   Airport.Name                         52704 non-null  object
10  Injury.Severity                      87889 non-null  object
11  Aircraft.damage                      85695 non-null  object
12  Aircraft.Category                    32287 non-null  object
13  Registration.Number                 87507 non-null  object
14  Make                                88826 non-null  object
15  Model                               88797 non-null  object
16  Amateur.Built                       88787 non-null  object
17  Number.of.Engines                   82805 non-null  float64
18  Engine.Type                         81793 non-null  object
19  FAR.Description                     32023 non-null  object
20  Schedule                            12582 non-null  object
21  Purpose.of.flight                   82697 non-null  object
22  Air.carrier                         16648 non-null  object
23  Total.Fatal.Injuries                77488 non-null  float64
24  Total.Serious.Injuries              76379 non-null  float64
25  Total.Minor.Injuries                76956 non-null  float64
26  Total.Uninjured                     82977 non-null  float64
27  Weather.Condition                   84397 non-null  object
28  Broad.phase.of.flight               61724 non-null  object
29  Report.Status                       82505 non-null  object
30  Publication.Date                    75118 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.0+ MB
```

In [34]: `df.describe()` *#Statistical summary of my data frame*

Out[34]:

	Number.of.Engines	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries
count	82805.000000	77488.000000	76379.000000	76956.000000
mean	1.146585	0.647855	0.279881	0.357061
std	0.446510	5.485960	1.544084	2.235625
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000	0.000000
max	8.000000	349.000000	161.000000	380.000000

```
In [35]: df.isnull().sum() #Check for null values
```

```
Out[35]: Event.Id          0
Investigation.Type        0
Accident.Number           0
Event.Date                0
Location                  52
Country                   226
Latitude                  54507
Longitude                 54516
Airport.Code              38757
Airport.Name              36185
Injury.Severity           1000
Aircraft.damage           3194
Aircraft.Category         56602
Registration.Number       1382
Make                       63
Model                     92
Amateur.Built             102
Number.of.Engines         6084
Engine.Type               7096
FAR.Description           56866
Schedule                  76307
Purpose.of.flight         6192
Air.carrier               72241
Total.Fatal.Injuries      11401
Total.Serious.Injuries    12510
Total.Minor.Injuries      11933
Total.Uninjured           5912
Weather.Condition         4492
Broad.phase.of.flight     27165
Report.Status             6384
Publication.Date          13771
dtype: int64
```

DATA PREPARATION

Data Cleaning

Data cleaning and wrangling is essential to the data science process. I do these by first checking and dropping all the duplicates in the data frame.


Next I standardize all the column names by removing white space, using underscore as a separator and making all the characters lower case making them easier to call.

I then filter out all the data older than 6 years because of constant innovation and changing technologies in the aviation space


Dropping columns with more than half null values

I impute all numerical columns with zero[0] and all the categorical with "unknown"

Lastly I drop all the non essential columns and finally I drop all rows with null values

```
In [36]:  #Check and drop duplicates  
df.duplicated().sum()
```

Out[36]: 0

```
In [37]:  #Standardize column names by removing white space, lowercase, using underscores  
df.columns=df.columns.str.strip().str.lower().str.replace(" ", "_").str.strip()  
  
df.columns
```

Out[37]: Index(['event_id', 'investigation_type', 'accident_number', 'event_date',
 'location', 'country', 'latitude', 'longitude', 'airport_code',
 'airport_name', 'injury_severity', 'aircraft_damage',
 'aircraft_category', 'registration_number', 'make', 'model',
 'amateur_built', 'number_of_engines', 'engine_type', 'far_description',
 'schedule', 'purpose_of_flight', 'air_carrier', 'total_fatal_injuries',
 'total_serious_injuries', 'total_minor_injuries', 'total_uninjured',
 'weather_condition', 'broad_phase_of_flight', 'report_status',
 'publication_date'],
 dtype='object')

```
In [38]: #Drop data that is more than 6 years old

#Covert event_date to date time format
df["event_date"]=pd.to_datetime(df["event_date"])
df["event_date"].isnull().sum() #check for any null values of thr event

start="2017-01-01" #Define the start date

df=df[df["event_date"]>=start] #Filter the data set

df.head(10) #Check whether its been correctly filtered
```

Out[38]:

	event_id	investigation_type	accident_number	event_date	location
79402	20170110X14448	Accident	GAA17CA108	2017-01-01	Casa Grande, AZ
79403	20170117X42903	Accident	CEN17LA079	2017-01-01	New Braunfels, TX
79404	20170103X14851	Accident	WPR17FA045	2017-01-02	Payson, AZ
79405	20170103X43747	Accident	WPR17LA046	2017-01-03	Paradise, MT
79406	20170105X45917	Accident	WPR17LA048	2017-01-04	Napa, CA
79407	20170105X51011	Accident	CEN17LA068	2017-01-04	Nacogdoches, TX
79408	20170105X81632	Accident	CEN17FA067	2017-01-04	Brookfield, WI
79409	20170105X41106	Accident	WPR17FA047	2017-01-05	San Pedro, CA
79410	20170105X74726	Accident	CEN17FA071	2017-01-05	Gurdon, AR
79411	20170105X85226	Accident	ERA17LA078	2017-01-05	Atlanta, GA

In [39]: `df.shape` *#The new shape of the filtered data*

Out[39]: (9487, 31)

In [40]: `df.isnull().sum()` *#Check the null values of the new data*

Out[40]:

event_id	0
investigation_type	0
accident_number	0
event_date	0
location	0
country	0
latitude	1347
longitude	1347
airport_code	3979
airport_name	4162
injury_severity	559
aircraft_damage	780
aircraft_category	192
registration_number	2
make	0
model	0
amateur_built	0
number_of_engines	2019
engine_type	3795
for_description	100

```
In [41]: #Drop columns with more than 50% of missing values  
thresh=len(df)*0.5  
df=df.loc[:,df.isnull().sum()<=thresh]  
  
df.isnull().sum()
```

```
Out[41]: event_id          0  
investigation_type      0  
accident_number         0  
event_date              0  
location                0  
country                 0  
latitude               1347  
longitude              1347  
airport_code           3979  
airport_name           4162  
injury_severity         559  
aircraft_damage         780  
aircraft_category       192  
registration_number      2  
make                    0  
model                   0  
amateur_built           0  
number_of_engines       2019  
engine_type             3795  
far_description          199  
purpose_of_flight       2194  
total_fatal_injuries     0  
total_serious_injuries   0  
total_minor_injuries     0  
total_uninjured          0  
weather_condition       2320  
report_status           4522  
publication_date         953  
dtype: int64
```



```
In [42]: #Filling in numerical columns  
numcol=["total_fatal_injuries", "total_serious_injuries","total_minor_i  
#The numerical columns will be filled with [0]  
df[numcol]=df[numcol].fillna(0)  
df.isnull().sum()
```

```
Out[42]: event_id          0  
investigation_type      0  
accident_number         0  
event_date              0  
location                0  
country                 0  
latitude               1347  
longitude               1347  
airport_code            3979  
airport_name            4162  
injury_severity         559  
aircraft_damage         780  
aircraft_category       192  
registration_number      2  
make                    0  
model                   0  
amateur_built            0  
number_of_engines        0  
engine_type             3795  
far_description          199  
purpose_of_flight       2194  
total_fatal_injuries      0  
total_serious_injuries    0  
total_minor_injuries      0  
total_uninjured          0  
weather_condition       2320  
report_status            4522  
publication_date         953  
dtype: int64
```

```
In [43]: #Categorical columns  
catcol=["location","injury_severity","weather_condition","country","eng  
#They will be filled with unknown  
df[catcol]=df[catcol].fillna("Unk")  
  
df.isnull().sum()
```

```
Out[43]: event_id          0  
investigation_type      0  
accident_number        0  
event_date             0  
location               0  
country               0  
latitude              1347  
longitude             1347  
airport_code          3979  
airport_name          4162  
injury_severity        0  
aircraft_damage        0  
aircraft_category      0  
registration_number     2  
make                  0  
model                 0  
amateur_built          0  
number_of_engines      0  
engine_type           0  
far_description        199  
purpose_of_flight     2194  
total_fatal_injuries    0  
total_serious_injuries  0  
total_minor_injuries   0  
total_uninjured        0  
weather_condition      0  
report_status          4522  
publication_date       953  
dtype: int64
```

```
In [44]: #Dropping columns not useful for the analysis"fl  
df=df.drop(columns=["airport_code","airport_name","registration_number"]
```

```
In [45]: df.isnull().sum()
```

```
Out[45]: event_id          0
investigation_type      0
event_date             0
location              0
country               0
injury_severity        0
aircraft_damage        0
aircraft_category      0
make                  0
model                 0
amateur_built          0
number_of_engines      0
engine_type            0
total_fatal_injuries    0
total_serious_injuries  0
total_minor_injuries    0
total_uninjured        0
weather_condition      0
dtype: int64
```

```
In [46]: #Drop all other null values
df=df.dropna()
```

```
In [47]: df.shape
```

```
Out[47]: (9487, 18)
```

```
In [48]: #Check our final data set with no null values
df.isnull().sum()
```

```
Out[48]: event_id          0
investigation_type      0
event_date             0
location              0
country               0
injury_severity        0
aircraft_damage        0
aircraft_category      0
make                  0
model                 0
amateur_built          0
number_of_engines      0
engine_type            0
total_fatal_injuries    0
total_serious_injuries  0
total_minor_injuries    0
total_uninjured        0
weather_condition      0
dtype: int64
```

In [49]:

```
# Convert all names in the 'make' column to lowercase
df['make'] = df['make'].str.lower().str.replace(" ", "_")

# Display unique makes to confirm the changes
unique_makes = df['make'].unique()

# Print the first 50 unique makes and total count
print("Unique Aircraft Makes in Lowercase:")
print(unique_makes[:50]) # Display the first 50 unique makes
print(f"Total unique aircraft makes: {len(unique_makes)}")
```

```
Unique Aircraft Makes in Lowercase:
['larson_roger_h' 'quicksilver' 'cessna' 'softex_invest_llc'
 'schosanski_john_h' 'mooney_aircraft_corp.' 'demmer'
 'robinson_helicopter' 'columbia_aircraft_mfg'
 'aircraft_mfg_&_development_co' 'boeing' 'canadair' 'cirrus'
 'christen_industries_inc' 'boeing_company' 'maule' 'fields' 'mooney'
 'piper' 'sharpe_william_l' 'beech' 'piper_aircraft_corporation' 'aer
onca'
 'american_champion_aircraft' 'textron_aviation_inc'
 'american_legend_aircraft_co' 'johnson' 'air_tractor_inc'
 'hawker_beechcraft' 'avro' 'hawker_beechcraft_corp' 'bellanca' 'robi
nson'
 'bombardier_inc' 'bell' 'cirrus_design_corp'
 'grumman_american_corporation' 'eurocopter_deutschland_gmbh' 'ercoup
e'
 'great_lakes' 'vaughan_gerald_r' 'aviat_aircraft_inc' 'zenith' 'coo
k'
 'robinson_michael_e' 's_c_aerostar_s_a' 'bell_helicopter_textron'
 'shell_john' 'saab' 'md_helicopter']
Total unique aircraft makes: 1405
```

In [50]: df.shape

Out[50]: (9487, 18)

In [51]: df.to_csv(r"C:\Users\User\Desktop\Flight-Data\AviationData_clean.csv", i

ANALYSIS

1. Analyse injury severity by weather condition

We can see that majority of accidents, even though they are non-fatal, happen during clear skies and weather conditions when pilots rely on visual references and stimuli

In [52]:

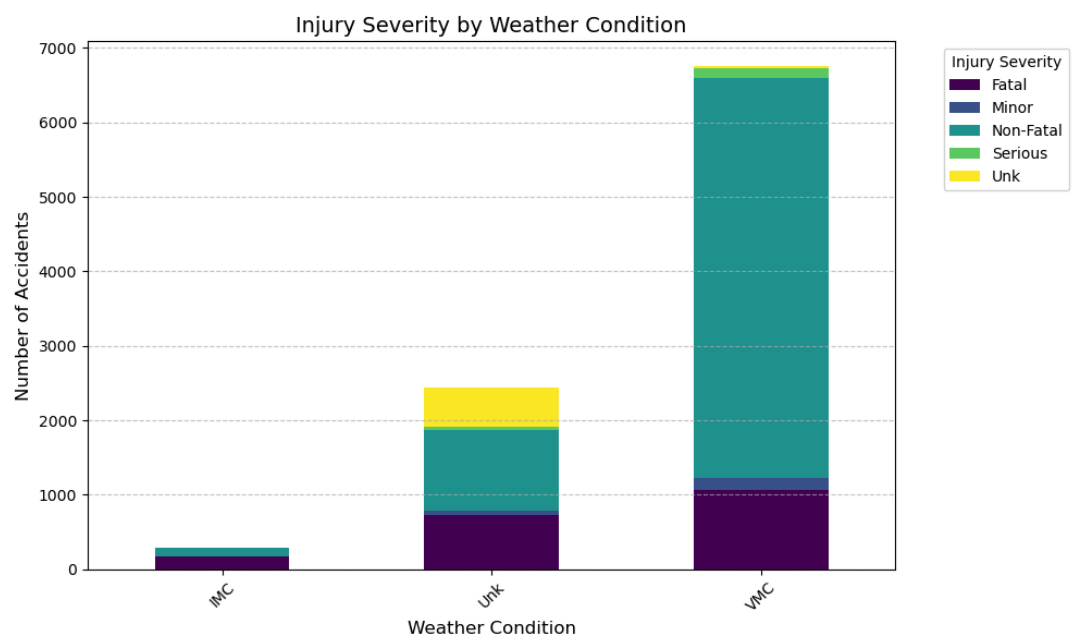
```

# Cross-tabulation of weather condition and injury severity
weather_injury = pd.crosstab(df['weather_condition'], df['injury_severity'])

# Visualize the relationship
weather_injury.plot(kind='bar', stacked=True, figsize=(10, 6), colormap='magma')
plt.title('Injury Severity by Weather Condition', fontsize=14)
plt.xlabel('Weather Condition', fontsize=12)
plt.ylabel('Number of Accidents', fontsize=12)
plt.xticks(rotation=45)
plt.legend(title='Injury Severity', bbox_to_anchor=(1.05, 1), loc='upper right')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()

plt.show()

```



2. Aircraft make & model by fatalities per accidents

Modern aircrafts are all relatively safe due to strict regulations by relevant authorities. Even though we can clearly see the Airbus A320 and Boeing 737 have more accidents than the rest



```
In [53]: # Aggregate accident and fatality data by make and model
safest_models = df.groupby(['make', 'model']).agg({
    'event_id': 'count', # Total accidents
    'total_fatal_injuries': 'sum' # Total fatalities
}).reset_index()

safest_models=safest_models.sort_values(by="total_fatal_injuries", ascending=True)
# safest_models.tail(50)
# "embraer" in safest_models["make"].unique()

# biggest Commercial Plane manufacturer
commercial_planes=safest_models[safest_models["make"].str.lower().isin(['embraer', 'boeing', 'airbus'])]
commercial_planes.head()

# Get the top 10 models with the least fatalities
top_10 = commercial_planes.head(10)


# Get the bottom 10 models with the most fatalities
bottom_10 = commercial_planes.tail(10)

# Plot side-by-side bar charts
plt.figure(figsize=(12, 6))

# Plot top 10 models with the least fatalities
plt.subplot(1, 2, 1)
plt.bar(top_10['model'], top_10['total_fatal_injuries'], color='green')
plt.title('Top 10 Safest Models (Least Fatal Injuries)', fontsize=12)
plt.xlabel('Aircraft Model', fontsize=10)
plt.ylabel('Total Fatal Injuries', fontsize=10)
plt.xticks(rotation=45)

# Plot bottom 10 models with the most fatalities
plt.subplot(1, 2, 2)
plt.bar(bottom_10['model'], bottom_10['total_fatal_injuries'], color='red')
plt.title('Bottom 10 Models (Most Fatal Injuries)', fontsize=12)
plt.xlabel('Aircraft Model', fontsize=10)
plt.xticks(rotation=45)

# Adjust layout for readability
plt.tight_layout()
plt.show()
```




3. Analyse total number of accidents by weather conditions

This graph shows that majority of accidents happen during clear weather conditions while pilots use visual flight rules(VFR)

In [54]:

```
# Count accidents by each weather condition
weather_accidents = df['weather_condition'].value_counts()

# Plot weather accidents
plt.figure(figsize=(10, 6))
weather_accidents.plot(kind='bar', color='skyblue')
plt.title('Accidents by Weather Condition', fontsize=14)
plt.xlabel('Weather Condition', fontsize=12)
plt.ylabel('Number of Accidents', fontsize=12)
plt.xticks(rotation=45)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



CONCLUSION

The analysis conducted leads to the following observations and recommendations

1. Invest in advanced training programs for pilots to enhance navigational skills and advanced situational awareness.

This is because most accidents happen during clear skies/ visual meteorological conditions suggesting that over reliance in visual references leads to errors.

2. Consider other factors such as RELIABILITY, EFFICIENCY, VERSATILITY AND PRICE while purchasing aircrafts instead of accident data.

This is because even though we can clearly see the Airbus A320 and Boeing 737 have more accidents than the rest it's crucial to note that it's not a reflection on poor safety. These two are the most widely used planes

3. Adopt advanced navigation and auto pilot systems to enhance safety even clear weather conditions

Most accidents happen during VFR conditions which are deceptively safer but pose unique challenges